

## Sparkify Project Workspace

This workspace contains a tiny subset (128MB) of the full dataset available (12GB). Feel free to use this workspace to build your project, or to explore a smaller subset with Spark before deploying your cluster on the cloud. Instructions for setting up your Spark cluster is included in the last lesson of the Extracurricular Spark Course content.

You can follow the steps below to guide your data analysis and model building portion of this project.

```
In [111]: # import libraries
import numpy as np
import pandas as pd
import time
import datetime
import seaborn as sns
import matplotlib.pyplot as plt
from statsmodels.stats.proportion import proportions_ztest

%matplotlib inline

from pyspark.sql import SparkSession, Window
from pyspark.sql.functions import avg, stddev, split, udf, isnull, first
, col, format_number, rand
from pyspark.sql.functions import min as fmin
from pyspark.sql.functions import max as fmax
from pyspark.sql.types import IntegerType, FloatType

import re
import copy
import seaborn as sns
import matplotlib.pyplot as plt
from statsmodels.stats.proportion import proportions_ztest

%matplotlib inline

from pyspark.ml import Pipeline
from pyspark.ml.classification import LogisticRegression, RandomForestCl
assifier, GBTCClassifier, DecisionTreeClassifier, NaiveBayes
from pyspark.ml.feature import StandardScaler, VectorAssembler
from pyspark.ml.evaluation import MulticlassClassificationEvaluator
from pyspark.mllib.evaluation import MulticlassMetrics
from pyspark.ml.tuning import CrossValidator, ParamGridBuilder

import re
import copy

import seaborn as sns
import matplotlib.pyplot as plt
from statsmodels.stats.proportion import proportions_ztest

%matplotlib inline
```

```
In [8]: sns.set(style="whitegrid")
```

```
In [9]: # create a Spark session
spark_1 = SparkSession.builder.appName('Sparkify_local').getOrCreate()
```

# Load and Clean Dataset

In this workspace, the mini-dataset file is `mini_sparkify_event_data.json` . Load and clean the dataset, checking for invalid or missing data - for example, records without `userids` or `sessionids`.

## Reading the data

```
In [10]: df1 = spark_1.read.json('mini_sparkify_event_data.json')
df1.show(1)
```

artist	auth	firstName	gender	itemInSession	lastName	length	level	location	method	page	registration	sessionId	song	status	ts	userAgent	userId
Martha Tilston	Logged In	Colin	M	50	Freeman	277.89016	paid	Bakersfield, CA	PUT	NextSong	1538173362000	29	Rock pools	200	1538352117000	Mozilla/5.0 (Wind...	30

only showing top 1 row

```
In [11]: df1.printSchema()

root
|-- artist: string (nullable = true)
|-- auth: string (nullable = true)
|-- firstName: string (nullable = true)
|-- gender: string (nullable = true)
|-- itemInSession: long (nullable = true)
|-- lastName: string (nullable = true)
|-- length: double (nullable = true)
|-- level: string (nullable = true)
|-- location: string (nullable = true)
|-- method: string (nullable = true)
|-- page: string (nullable = true)
|-- registration: long (nullable = true)
|-- sessionId: long (nullable = true)
|-- song: string (nullable = true)
|-- status: long (nullable = true)
|-- ts: long (nullable = true)
|-- userAgent: string (nullable = true)
|-- userId: string (nullable = true)
```

```
In [12]: df1.count()
```

```
Out[12]: 286500
```

```
In [13]: df1.describe('userId').show()
```

```
+-----+-----+
|summary|      userId|
+-----+-----+
|   count|      286500|
|   mean|59682.02278593872|
| stddev|109091.9499991047|
|    min|              |
|    max|              99|
+-----+-----+
```

```
In [14]: df1.describe('sessionId').show()
```

```
+-----+-----+
|summary|      sessionId|
+-----+-----+
|   count|      286500|
|   mean|1041.526554973822|
| stddev|726.7762634630741|
|    min|              1|
|    max|          2474|
+-----+-----+
```

```
In [15]: df1.persist()
```

```
Out[15]: DataFrame[artist: string, auth: string, firstName: string, gender: string, itemInSession: bigint, lastName: string, length: double, level: string, location: string, method: string, page: string, registration: bigint, sessionId: bigint, song: string, status: bigint, ts: bigint, userAgent: string, userId: string]
```

to aid our analysis, we will need to clean our data so it is more consistent for our use.

## Cleaning the Data

As the first step we will remove the missing values and NaNs

```
In [16]: df1.filter(isnull(df1['sessionId'])).count()
```

```
Out[16]: 0
```

```
In [17]: df1.filter(isnull(df1['userId'])).count()
```

```
Out[17]: 0
```

```
In [18]: df1.filter(df1['userId']=='').count()
```

```
Out[18]: 8346
```

```
In [19]: df1.filter(df1['sessionId']=='').count()
```

```
Out[19]: 0
```

```
In [20]: df1_clean = df1.filter(df1['userId']!='')
```

```
In [21]: ''' One noticeable thing we see in the data given to us is
that the time is not in a readable format so we can try and get it
into the form that is human readable
'''

time_old = udf(lambda x: datetime.datetime.fromtimestamp(x / 1000.0).strftime("%Y-%m-%d %H:%M:%S"))

df_clean = df1_clean.withColumn("time", time_old(df1_clean.ts))
```

```
In [22]: df_clean.take(1)
```

```
Out[22]: [Row(artist='Martha Tilston', auth='Logged In', firstName='Colin', gender='M', itemInSession=50, lastName='Freeman', length=277.89016, level='paid', location='Bakersfield, CA', method='PUT', page='NextSong', registration=1538173362000, sessionId=29, song='Rockpools', status=200, ts=1538352117000, userAgent='Mozilla/5.0 (Windows NT 6.1; WOW64; rv:31.0) Gecko/20100101 Firefox/31.0', userId='30', time='2018-10-01 00:01:57')]
```

## Exploratory Data Analysis

When you're working with the full dataset, perform EDA by loading a small subset of the data and doing basic manipulations within Spark. In this workspace, you are already provided a small subset of data you can explore.

```
In [23]: # lets explore the unique values of the 'page' column
df_clean.select("page").dropDuplicates().show()
```

```
+-----+
|           page|
+-----+
|           Cancel|
| Submit Downgrade|
|       Thumbs Down|
|             Home|
|         Downgrade|
|       Roll Advert|
|           Logout|
|       Save Settings|
| Cancellation Conf...|
|             About|
|           Settings|
| Add to Playlist|
|       Add Friend|
|       NextSong|
|       Thumbs Up|
|             Help|
|           Upgrade|
|           Error|
| Submit Upgrade|
+-----+
```

```
In [24]: df_clean.filter(df_clean.page=="Cancellation Confirmation").select("userID").dropDuplicates().show(15)
```

```
+-----+
|userID|
+-----+
|    125|
|     51|
|     54|
| 100014|
|     101|
|     29|
| 100021|
|     87|
|     73|
|      3|
|     28|
| 100022|
| 100025|
| 300007|
| 100006|
+-----+
```

only showing top 15 rows

```
In [25]: df_clean.select(["userId", "page", "time", "level", "song", "sessionId"])\n         .where(df_clean.userId == "30").sort("time").show(100)
```

userId	page	time	level	song
sessionId				
29	30	NextSong 2018-10-01 00:01:57	paid	Rockpools
29	30	NextSong 2018-10-01 00:06:34	paid	Time For Miracles
29	30	NextSong 2018-10-01 00:11:16	paid	Harder Better Fas...
29	30	NextSong 2018-10-01 00:14:59	paid	Passengers (Old A...
29	30	Add to Playlist 2018-10-01 00:15:05	paid	null
29	30	NextSong 2018-10-01 00:18:04	paid	Fuck Kitty
29	30	NextSong 2018-10-01 00:20:18	paid	Jade
29	30	NextSong 2018-10-01 00:24:01	paid	So-Called Friends
29	30	NextSong 2018-10-01 00:28:07	paid	Represent
29	30	NextSong 2018-10-01 00:31:49	paid	Here I Am
29	30	NextSong 2018-10-01 00:35:32	paid	Rebirthing (Album...
29	30	NextSong 2018-10-01 00:39:25	paid	Dog Days Are Over...
29	30	NextSong 2018-10-01 00:43:04	paid	Tomorrow Is A Lon...
29	30	NextSong 2018-10-01 00:46:46	paid	Halloween Spooks
29	30	NextSong 2018-10-01 00:49:05	paid	Stronger
29	30	NextSong 2018-10-01 00:54:16	paid	Dis Iz Brick City
29	30	NextSong 2018-10-01 00:57:53	paid	Move Along
29	30	NextSong 2018-10-01 01:01:51	paid	Manhattan
29	30	NextSong 2018-10-01 01:05:15	paid	Undo
29	30	NextSong 2018-10-01 01:11:03	paid	The Big Gundown
29	30	NextSong 2018-10-01 01:15:23	paid	Black Bird
29	30	Thumbs Down 2018-10-01 01:15:24	paid	null
29	30	NextSong 2018-10-01 01:18:27	paid	Nausea
29	30	NextSong 2018-10-01 01:21:43	paid	Matricide
29	30	NextSong 2018-10-01 01:27:03	paid	Valerie
29	30	NextSong 2018-10-01 01:30:52	paid	Margarita



```

29 |
| 30 |      NextSong|2018-10-01 01:34:08| paid|      Le Jardin d'Hiver|
29 |
| 30 |      Thumbs Up|2018-10-01 01:34:09| paid|                        null|
29 |
| 30 |      NextSong|2018-10-01 01:39:50| paid|Soon As I Get Hom...|
29 |
| 30 |      Thumbs Up|2018-10-01 01:39:51| paid|                        null|
29 |
| 30 |      NextSong|2018-10-01 01:45:14| paid|      Vamos a la Playa|
29 |
| 30 |      NextSong|2018-10-01 01:48:52| paid|                        Perfecta|
29 |
| 30 |      NextSong|2018-10-01 01:52:41| paid|                        Requiem|
29 |
| 30 |      NextSong|2018-10-01 01:57:34| paid|      Who Can Compare|
29 |
| 30 |      NextSong|2018-10-01 02:02:00| paid|240 Years Before ...|
29 |
| 30 |      NextSong|2018-10-01 02:25:20| paid|      Rosa Pastel|
29 |
| 30 |      Roll Advert|2018-10-01 02:25:37| paid|                        null|
29 |
| 30 |      NextSong|2018-10-01 02:28:25| paid|I'm Ready (Album ...|
29 |
| 30 |      NextSong|2018-10-01 02:31:47| paid|      No Other Saviour|
29 |
| 30 |      NextSong|2018-10-01 02:36:12| paid|                        Hints|
29 |
| 30 |      NextSong|2018-10-01 02:39:00| paid|                        Yellow|
29 |
| 30 |Add to Playlist|2018-10-01 02:39:35| paid|                        null|
29 |
| 30 |      NextSong|2018-10-06 07:23:50| paid|Kill The Director...|
264 |
| 30 |      Add Friend|2018-10-06 07:23:51| paid|                        null|
264 |
| 30 |      Home|2018-10-07 21:37:06| paid|                        null|
532 |
| 30 |      NextSong|2018-10-07 21:37:20| paid|      Third Party|
532 |
| 30 |      NextSong|2018-10-07 21:43:11| paid|      It's Working|
532 |
| 30 |      NextSong|2018-10-07 21:47:17| paid|I'll Remember Apr...|
532 |
| 30 |      NextSong|2018-10-07 21:50:34| paid|      Atmosphere Station|
532 |
| 30 |      NextSong|2018-10-07 21:53:45| paid|      There_ There|
532 |
| 30 |      NextSong|2018-10-07 21:59:08| paid|'Til We Die (Albu...|
532 |
| 30 |      NextSong|2018-10-07 22:04:53| paid|Lies (Album Version)|
532 |
| 30 |      NextSong|2018-10-07 22:07:52| paid|The Ballad of Mic...|
532 |
| 30 |      NextSong|2018-10-07 22:11:42| paid|                        Gears|
532 |

```

532		30		NextSong		2018-10-07 22:15:41		paid		These Eyes	
532		30		Help		2018-10-07 22:15:49		paid		null	
532		30		Home		2018-10-07 22:16:33		paid		null	
532		30		NextSong		2018-10-07 22:20:12		paid		Stronger	
532		30		NextSong		2018-10-07 22:25:23		paid		This Is Such A Pity	
532		30		NextSong		2018-10-07 22:28:47		paid		Strut (1993 Digit...	
532		30		NextSong		2018-10-07 22:32:46		paid		Fader	
532		30		NextSong		2018-10-07 22:35:58		paid		Se Quiere_ Se Mata	
532		30		NextSong		2018-10-07 22:39:36		paid		Somewhere	
532		30		NextSong		2018-10-07 22:42:25		paid		Canada	
532		30		NextSong		2018-10-07 22:46:21		paid		Dream On Dreamer	
532		30		NextSong		2018-10-07 22:50:21		paid		HÃ   Âi Uma MÃ   Âºsi...	
532		532		NextSong		2018-10-07 22:53:57		paid		Money Ain't A Thang	
532		30		NextSong		2018-10-07 22:58:11		paid		Howlin For You	
532		30		NextSong		2018-10-07 23:01:22		paid		Fireflies	
532		30		NextSong		2018-10-07 23:05:07		paid		Stadium Love	
532		30		Downgrade		2018-10-07 23:05:26		paid		null	
532		30		NextSong		2018-10-07 23:09:19		paid		Again & Again	
532		30		NextSong		2018-10-07 23:12:04		paid		The Trooper (1998...	
532		30		NextSong		2018-10-07 23:16:16		paid		Get Off My Elevator	
532		30		NextSong		2018-10-07 23:18:34		paid		Speed Trials	
532		30		NextSong		2018-10-07 23:21:35		paid		Invalid	
532		30		NextSong		2018-10-07 23:25:28		paid		Frenchy s	
532		30		NextSong		2018-10-07 23:28:22		paid		Between Two Lungs	
532		30		NextSong		2018-10-07 23:32:31		paid		Hiding	
532		30		NextSong		2018-10-07 23:36:17		paid		Alone Again (Natu...	
532		30		NextSong		2018-10-07 23:39:55		paid		You Found Me	
532		30		NextSong		2018-10-07 23:43:55		paid		Our Song	
532		30		NextSong		2018-10-07 23:47:16		paid		Rock Hard	

```

532|
| 30|      NextSong|2018-10-07 23:50:30| paid| Waiting For Tonight|
532|
| 30|      NextSong|2018-10-07 23:54:35| paid| Pioneer to The Falls|
532|
| 30|      NextSong|2018-10-08 00:00:16| paid|      You're The One|
532|
| 30|      NextSong|2018-10-08 00:04:15| paid|      Trash And Ready|
532|
| 30|      NextSong|2018-10-08 00:07:08| paid|      Keep On Movin'|
532|
| 30|      NextSong|2018-10-08 00:13:09| paid|      Quase Um Segundo|
532|
| 30|      NextSong|2018-10-08 00:16:03| paid| Costumbres Argent...|
532|
| 30|      NextSong|2018-10-08 00:19:20| paid| Nine Times Out Of...|
532|
| 30|      NextSong|2018-10-08 00:21:30| paid|      Bass Solo|
532|
| 30|      NextSong|2018-10-08 00:26:27| paid| The Memory Remains|
532|
| 30|      NextSong|2018-10-08 00:31:06| paid|      Clock|
532|
| 30|      NextSong|2018-10-08 00:35:23| paid| I Don t Wanna Go ...|
532|
| 30|      NextSong|2018-10-08 00:37:07| paid|      She's So Cold|
532|
| 30|      NextSong|2018-10-08 00:41:20| paid| Drops In The River|
532|
| 30|      NextSong|2018-10-08 00:45:33| paid|      Didgeridoo|
532|
| 30|      NextSong|2018-10-08 00:52:43| paid| Sex_ Love & Money|
532|
| 30|      NextSong|2018-10-08 00:56:53| paid|      Rayando el sol|
532|
+-----+-----+-----+-----+-----+
-----+
only showing top 100 rows

```

In [26]: *# one way to go around is using Spark SQL to create temporary view and perform querying*

```
df_clean.createOrReplaceTempView("sparkify_clean")
```

In [27]: *# lets take a peek into different types of authorization*

```
spark_1.sql(''
            SELECT DISTINCT(auth)
            FROM sparkify_clean
            '').show()
```

```
+-----+
|      auth|
+-----+
|Cancelled|
|Logged In|
+-----+
```

In [29]: spark = spark\_1

In [30]: *# now let us look for user count in each category*

```
spark.sql(''
          SELECT auth,COUNT(DISTINCT userId) AS user_counts
          FROM sparkify_clean
          GROUP BY auth
          ORDER BY user_counts DESC
          '').show()
```

```
+-----+-----+
|      auth|user_counts|
+-----+-----+
|Logged In|         225|
|Cancelled|          52|
+-----+-----+
```


In [31]: *# One possible avenue of exploration is how is the gender distributed in the data*

```
gender_data = spark.sql(''
                        SELECT gender,COUNT(DISTINCT userId) AS user_counts
                        FROM sparkify_clean
                        GROUP BY gender
                        ORDER BY user_counts DESC
                        '')
gender_data.show()
```

```
+-----+-----+
|gender|user_counts|
+-----+-----+
|      M|         121|
|      F|         104|
+-----+-----+
```

We can see that our data consists of 121 Males and 104 Females. Let us try and visualize it.


```
sns.barplot(x='gender',y='user_counts',data=gender_data.toPandas());
```




```
# Let us look into how the data is divided by location
location_data = spark.sql('''
    SELECT location,COUNT(DISTINCT userId) AS user_counts
    FROM sparkify_clean
    GROUP BY location
    ORDER BY user_counts DESC
''').toPandas()
```

```
#split city and state
location_data = location_data.join(location_data['location'].str.split(
    ',', expand=True).rename(columns={0: 'city', 1: 'state'})).drop('location', a
xis=1)
```

```
location_data.groupby('city')['user_counts'].sum().sort_values(ascending=False).plot(kind='bar',figsize=(17,5));
```




```
In [36]: location_data.groupby('state')['user_counts'].sum().sort_values(ascending=False).plot(kind='bar',figsize=(17,5));
```



```
In [37]: # One other way to look at data is from 'length' point of view
```

```
length_data = spark.sql('''
    SELECT Length
    FROM sparkify_clean
''')

sns.distplot(length_data.toPandas().dropna());
```



We can observe that mostly all of the lengths are between 0 - 500.

```
In [38]: # lets look at the data from level perspective

spark.sql('''
    SELECT level, count(Distinct userID) as usr_counts
    FROM sparkify_clean
    GROUP BY level
    ORDER BY usr_counts DESC
''').show()
```

```
+-----+-----+
|level|usr_counts|
+-----+-----+
| free|        195|
| paid|        165|
+-----+-----+
```

The data split shows that there are 2 distinct levels with 195 'free' and 165 'paid' accounts. We can also deduce that there are 135 users that have changed their subscription levels.

```
In [39]: # Data split by page

spark.sql('''
    select page, count(userID) as usr_counts
    from sparkify_clean
    group by page
    order by usr_counts DESC
''').show()
```

```
+-----+-----+
|           page|usr_counts|
+-----+-----+
|      NextSong|    228108|
|    Thumbs Up|    12551|
|        Home|    10082|
|Add to Playlist|    6526|
|    Add Friend|    4277|
|   Roll Advert|    3933|
|      Logout|    3226|
|Thumbs Down|    2546|
|   Downgrade|    2055|
|    Settings|    1514|
|      Help|    1454|
|    Upgrade|     499|
|    About|     495|
|  Save Settings|    310|
|      Error|    252|
|Submit Upgrade|    159|
|Submit Downgrade|     63|
|      Cancel|     52|
|Cancellation Conf...|     52|
+-----+-----+
```


```
In [45]: # Data cut by 'time'

time_data = spark.sql('''
    SELECT time,userId
    FROM sparkify_clean
''').toPandas()


time_data['time'] = pd.to_datetime(time_data['time'])

weekday_dict = {0:'Mon.',1:'Tues.',2:'Wed.',3:'Thur.',4:'Fri.',5:'Sat.',
6:'Sun.'}
time_data['weekday'] = time_data['time'].dt.weekday.map(weekday_dict)
time_data['day'] = time_data['time'].dt.day
time_data['hour'] = time_data['time'].dt.hour

time_data.groupby('weekday')['userId'].count().loc[list(weekday_dict.values())].plot(kind='bar');
```




```
In [46]: time_data.groupby('day')['userId'].count().plot(kind='bar');
```





```
In [47]: time_data.groupby('hour')['userId'].count().plot(kind='bar');
```



We can deduce that the user behave in some sort of periodic pattern. It can be seen that users prefer to use it late in the day like after 2 or 3 o'clock in noon which can be justified considering that people start getting bored of their jobs. Also it is observed that the user prefer to use sparkify on weekdays over weekends, which also can be explained as people like to spend time with friends and family over the weekends.

```
In [ ]:
```

```
In [41]: # Data cut by userAgent

userAgent_count = spark.sql('''
    select userAgent, count(Distinct userID) as usr_counts
    from sparkify_clean
    group by userAgent
    order by usr_counts DESC
''').toPandas()
```

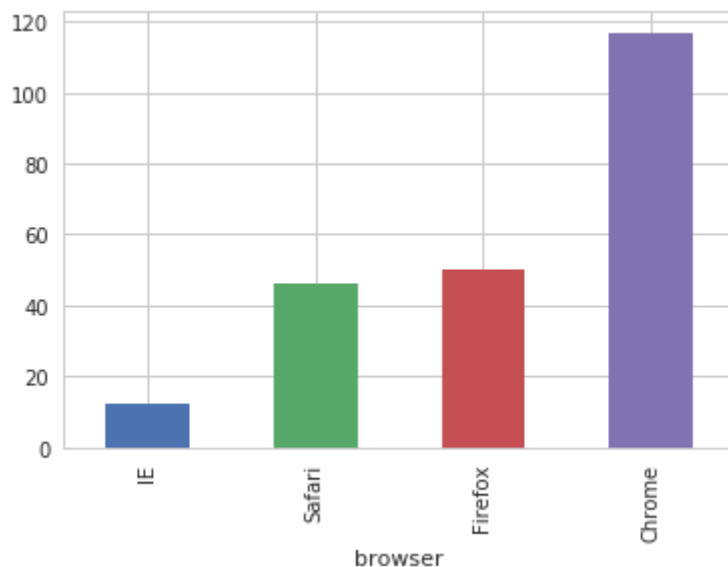
```
In [42]: def browser(x):
    if 'Firefox' in x:
        return 'Firefox'
    elif 'Safari' in x:
        if 'Chrome' in x:
            return 'Chrome'
        else:
            return 'Safari'
    elif 'Trident' in x:
        return 'IE'
    else:
        return np.NaN
```

```
In [43]: userAgent_count['browser'] = userAgent_count['userAgent'].apply(browser)


dict_platform = {'compatible': 'Windows', 'iPad': 'iPad', 'iPhone': 'i
Phone',
                 'Macintosh': 'Mac', 'Windows NT 5.1': 'Windows', 'Windows NT
6.0': 'Windows', 'Windows NT 6.1': 'Windows',
                 'Windows NT 6.2': 'Windows', 'Windows NT 6.3': 'Windows', 'X1
1': 'Linux'}

userAgent_count['platform'] = userAgent_count['userAgent'].str.extract(r
'\(([^\)]*)\)')[0].str.split(';').str[0].map(dict_platform)

userAgent_count.groupby('browser')['usr_counts'].sum().sort_values().plo
t(kind='bar');
```



```
In [44]: userAgent_count.groupby('platform')['usr_counts'].sum().sort_values().pl
ot(kind='bar');
```



In [ ]:

## Define Churn

Once you've done some preliminary analysis, create a column `Churn` to use as the label for your model. I suggest using the `Cancellation Confirmation` events to define your churn, which happen for both paid and free users. As a bonus task, you can also look into the `Downgrade` events.

## Explore Data

Once you've defined churn, perform some exploratory data analysis to observe the behavior for users who stayed vs users who churned. You can start by exploring aggregates on these two groups of users, observing how much of a specific action they experienced per a certain time unit or number of songs played.

```
In [48]: df_clean.select("page").dropDuplicates().show()
```

```
+-----+
|           page|
+-----+
|           Cancel|
| Submit Downgrade|
|       Thumbs Down|
|           Home|
|       Downgrade|
|       Roll Advert|
|         Logout|
|   Save Settings|
| Cancellation Conf...|
|           About|
|         Settings|
| Add to Playlist|
|       Add Friend|
|       NextSong|
|       Thumbs Up|
|           Help|
|         Upgrade|
|           Error|
|   Submit Upgrade|
+-----+
```

Lets us see for the customers who have selected the cancellation confirmation

```
In [50]: df_clean.filter(df_clean.page=="Cancellation Confirmation").select("userId").dropDuplicates().show(10)
```

```
+-----+  
|userId|  
+-----+  
|    125|  
|     51|  
|     54|  
|100014|  
|    101|  
|     29|  
|100021|  
|     87|  
|     73|  
|      3|  
+-----+
```

only showing top 10 rows

```
In [59]: # add time to see the time clear
get_time = udf(lambda x: datetime.datetime.fromtimestamp(x / 1000.0).strftime("%Y-%m-%d %H:%M:%S"))
df_clean = df_clean.withColumn("time", get_time(df_clean.ts))

df_clean.select(["userId", "page", "time", "level", "song", "sessionId"]).where(df_clean.userId == "87").sort("time").show()
```

userId	page	time	level	song	sessionId
86	87	Home	2018-10-01 06:47:15	free	null
86	87	NextSong	2018-10-01 06:49:00	free	Crotchopus
86	87	NextSong	2018-10-01 07:03:31	free	We Belong Together
86	87	NextSong	2018-10-01 07:07:58	free	Basic Space
86	87	NextSong	2018-10-01 07:11:06	free	Concertina
86	87	NextSong	2018-10-01 07:15:23	free	Crazy Little Thin...
86	87	NextSong	2018-10-01 07:18:33	free	Highway Song
86	87	Roll Advert	2018-10-01 07:18:44	free	null
86	87	NextSong	2018-10-01 07:29:05	free	Lipstick Traces (...)
86	87	NextSong	2018-10-01 07:31:31	free	Here's Your Revol...
86	87	NextSong	2018-10-01 07:35:22	free	Stack Shot Billy
86	87	NextSong	2018-10-01 07:38:43	free	I CAN'T GET STARTED
86	87	NextSong	2018-10-01 07:47:00	free	Queen Of Memphis ...
86	87	NextSong	2018-10-01 07:50:19	free	Up Up & Away
86	87	NextSong	2018-10-01 07:54:06	free	In One Ear
86	87	NextSong	2018-10-01 07:58:07	free	End Of The Road
86	87	NextSong	2018-10-01 08:01:26	free	Sex Out South
86	87	NextSong	2018-10-01 08:06:56	free	Don't Shake It Off
86	87	NextSong	2018-10-01 08:08:10	free	Karibien
273	87	NextSong	2018-10-01 15:12:31	free	Ã Â poca

only showing top 20 rows

In [60]: *# Adding churn column*

```
churn_usrs = df_clean.filter(df_clean.page=="Cancellation Confirmation")
.select("userId").dropDuplicates()
churn_usrs_list = [(row['userId']) for row in churn_usrs.collect()]
df_churn = df_clean.withColumn("churn", df_clean.userId.isin(churn_usrs_list))
```


In [61]: *# lets start to slice the data using this data which includes churn*

```
df_churn.dropDuplicates(["userId", "gender"]).groupby(["churn", "gender"]
).count().sort("churn").show()
```

```
+-----+-----+-----+
|churn|gender|count|
+-----+-----+-----+
|false|      M|    89|
|false|      F|    84|
| true|      F|    20|
| true|      M|    32|
+-----+-----+-----+
```

```
In [62]: df_pd = df_churn.dropDuplicates(["userId", "gender"]).groupby(["churn",
"gender"]).count().sort("churn").toPandas()
sns.barplot(x='churn', y='count', hue='gender', data=df_pd)
```

Out[62]: <matplotlib.axes.\_subplots.AxesSubplot at 0x7f2881ab5630>



```
In [63]: # 'page' use distribution between churn user and normal user


churn_count = df_churn.filter(df_churn.churn==1).count()
no_churn_count = df_churn.filter(df_churn.churn==0).count()

df_temp1 = df_churn.filter(df_churn.churn==1).filter(df_churn.page!="NextSong").groupby(["page"]).count()
df_temp1 = df_temp1.withColumn("cFreq", df_temp1["count"]/(churn_count/100)).sort("page")

df_temp2 = df_churn.filter(df_churn.churn==0).filter(df_churn.page!="NextSong").groupby(["page"]).count()
df_temp2 = df_temp2.withColumn("ncFreq", df_temp2["count"]/(no_churn_count/100)).sort("page")

df_pd = df_temp1.join(df_temp2, "page", "outer").drop("count").fillna(0)
df_pd.toPandas()
df_pd.plot.bar("page")
```

Out[63]: <matplotlib.axes.\_subplots.AxesSubplot at 0x7f288184fdd8>

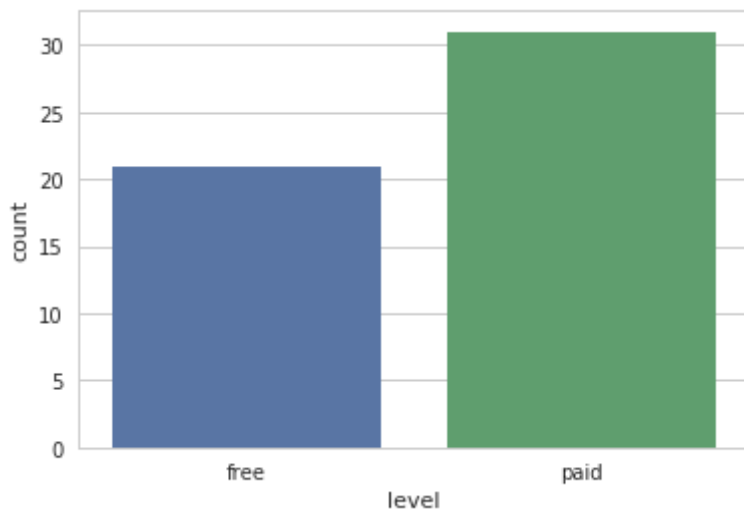




In [64]: *# Viewing the level when the user churned*

```
df_pd = df_churn.filter(df_churn.page=="Cancellation Confirmation").groupby("level").count().toPandas()  
sns.barplot(x="level", y="count", data=df_pd)
```


Out[64]: <matplotlib.axes.\_subplots.AxesSubplot at 0x7f2870d88c50>



In [67]: *# Slicing the data by distribution in operations in each session*

```
df_pd = df_churn.groupby("churn", "userId", "sessionId").count().toPandas()  
sns.boxplot(x='count', y='churn', orient="h", data=df_pd)
```

Out[67]: <matplotlib.axes.\_subplots.AxesSubplot at 0x7f28712aada0>



```
In [65]: # the time span of the log
```

```
df_churn.select("time").describe().show()
```

```
+-----+-----+
|summary|      time|
+-----+-----+
|   count|    278154|
|   mean|      null|
|  stddev|      null|
|    min|2018-10-01 00:01:57|
|    max|2018-12-03 01:11:16|
+-----+-----+
```

```
In [69]: # Time distribution in churn data



get_hour = udf(lambda x: datetime.datetime.fromtimestamp(x / 1000.0).hour)
df_churn = df_churn.withColumn("hour", get_hour(df_churn.ts))

get_weekday = udf(lambda x: datetime.datetime.fromtimestamp(x / 1000.0).strftime("%w"))
df_churn = df_churn.withColumn("weekday", get_weekday(df_churn.ts))


get_day = udf(lambda x: datetime.datetime.fromtimestamp(x / 1000.0).day)
df_churn = df_churn.withColumn("day", get_day(df_churn.ts))

def plot_cnt_by_churn(time):
    """
    This function use to plot the distribution of different dimension
    """
    df_pd = df_churn.filter(df_churn.page == "NextSong").groupby("churn", time).count().orderBy(df_churn[time].cast("float")).toPandas()
    df_pd[time] = pd.to_numeric(df_pd[time])
    df_pd[df_pd.churn==0].plot.bar(x=time, y='count', color='Blue', label='Not churn')
    df_pd[df_pd.churn==1].plot.bar(x=time, y='count', color='Red', label='Churn')



plot_cnt_by_churn("hour")
```



```
In [70]: plot_cnt_by_churn("weekday")
```



```
In [71]: plot_cnt_by_churn("day")
```



```
In [95]: df_churn.columns
```

```
Out[95]: ['artist',  
          'auth',  
          'firstName',  
          'gender',  
          'itemInSession',  
          'lastName',  
          'length',  
          'level',  
          'location',  
          'method',  
          'page',  
          'registration',  
          'sessionId',  
          'song',  
          'status',  
          'ts',  
          'userAgent',  
          'userId',  
          'time',  
          'churn',  
          'hour',  
          'weekday',  
          'day']
```

```
In [94]: df_churn.head(5)
```

```
Out[94]: [Row(artist='Martha Tilston', auth='Logged In', firstName='Colin', gender='M', itemInSession=50, lastName='Freeman', length=277.89016, level='paid', location='Bakersfield, CA', method='PUT', page='NextSong', registration=1538173362000, sessionId=29, song='Rockpools', status=200, ts=1538352117000, userAgent='Mozilla/5.0 (Windows NT 6.1; WOW64; rv:31.0) Gecko/20100101 Firefox/31.0', userId='30', time='2018-10-01 00:01:57', churn=False, hour='0', weekday='1', day='1'),
Row(artist='Five Iron Frenzy', auth='Logged In', firstName='Micah', gender='M', itemInSession=79, lastName='Long', length=236.09424, level='free', location='Boston-Cambridge-Newton, MA-NH', method='PUT', page='NextSong', registration=1538331630000, sessionId=8, song='Canada', status=200, ts=1538352180000, userAgent='"Mozilla/5.0 (Windows NT 6.1; WOW64) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/37.0.2062.103 Safari/537.36"', userId='9', time='2018-10-01 00:03:00', churn=False, hour='0', weekday='1', day='1'),
Row(artist='Adam Lambert', auth='Logged In', firstName='Colin', gender='M', itemInSession=51, lastName='Freeman', length=282.8273, level='paid', location='Bakersfield, CA', method='PUT', page='NextSong', registration=1538173362000, sessionId=29, song='Time For Miracles', status=200, ts=1538352394000, userAgent='Mozilla/5.0 (Windows NT 6.1; WOW64; rv:31.0) Gecko/20100101 Firefox/31.0', userId='30', time='2018-10-01 00:06:34', churn=False, hour='0', weekday='1', day='1'),
Row(artist='Enigma', auth='Logged In', firstName='Micah', gender='M', itemInSession=80, lastName='Long', length=262.71302, level='free', location='Boston-Cambridge-Newton, MA-NH', method='PUT', page='NextSong', registration=1538331630000, sessionId=8, song='Knocking On Forbidden Doors', status=200, ts=1538352416000, userAgent='"Mozilla/5.0 (Windows NT 6.1; WOW64) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/37.0.2062.103 Safari/537.36"', userId='9', time='2018-10-01 00:06:56', churn=False, hour='0', weekday='1', day='1'),
Row(artist='Daft Punk', auth='Logged In', firstName='Colin', gender='M', itemInSession=52, lastName='Freeman', length=223.60771, level='paid', location='Bakersfield, CA', method='PUT', page='NextSong', registration=1538173362000, sessionId=29, song='Harder Better Faster Stronger', status=200, ts=1538352676000, userAgent='Mozilla/5.0 (Windows NT 6.1; WOW64; rv:31.0) Gecko/20100101 Firefox/31.0', userId='30', time='2018-10-01 00:11:16', churn=False, hour='0', weekday='1', day='1')]
```

## Feature Engineering

Once you've familiarized yourself with the data, build out the features you find promising to train your model on. To work with the full dataset, you can follow the following steps.

- Write a script to extract the necessary features from the smaller subset of data
- Ensure that your script is scalable, using the best practices discussed in Lesson 3
- Try your script on the full data set, debugging your script if necessary

If you are working in the classroom workspace, you can just extract features based on the small subset of data contained here. Be sure to transfer over this work to the larger dataset when you work on your Spark cluster.



In [73]: *# Feature 1: number of days since registration*

```
user_max_ts = df_churn.groupby("userId").max("ts").sort("userId")
user_reg_ts = df_churn.select("userId", "registration").dropDuplicates()
                .sort("userId")
user_reg_days = user_reg_ts.join(user_max_ts, user_reg_ts.userId == user_max_ts.userId).select(user_reg_ts["userId"], ((user_max_ts["max(ts)"]-user_reg_ts["registration"])/(1000*60*60*24)).alias("regDay"))

user_reg_days.show(10)
```

```
+-----+-----+
|userId|      regDay|
+-----+-----+
|    10| 51.76265046296297|
|   100| 64.87377314814815|
|100001| 44.80021990740741|
|100002|160.47207175925925|
|100003|22.748113425925926|
|100004|172.44008101851853|
|100005| 85.19559027777778|
|100006| 9.127164351851851|
|100007|115.38761574074074|
|100008| 68.22856481481482|
+-----+-----+
```

only showing top 10 rows

In [85]: *# Feature 2: gender*

```
user_gender = df_churn.select("userId", "gender").dropDuplicates()
user_gender = user_gender.replace(["M", "F"], ["0", "1"], "gender")
user_gender = user_gender.select("userId", user_gender.gender.cast("int"))

user_gender.show(10)
```

```
+-----+-----+
|userId|gender|
+-----+-----+
|    44|    1|
|    46|    1|
|    41|    1|
|    72|    1|
|300023|    1|
|    39|    1|
|100010|    1|
|    40|    1|
|    94|    1|
|    35|    1|
+-----+-----+
```

only showing top 10 rows

In [77]: *# Feature 3: number of songs per session*

```
user_session_songs = df_churn.filter(df_churn.page=="NextSong").groupby(
    "userId", "sessionId").count()
user_session_songs_avg = user_session_songs.groupby("userId").agg(avg(us
er_session_songs["count"])).alias("avgSessionSongs").sort("userId")
```

```
user_session_songs_avg.show(10)
```

```
+-----+-----+
|userId| avgSessionSongs|
+-----+-----+
|      10| 112.16666666666667|
|     100|  78.88235294117646|
|  100001|           33.25|
|  100002|           48.75|
|  100003|           25.5|
|  100004|           47.1|
|  100005|           38.5|
|  100006|           26.0|
|  100007|           47.0|
|  100008| 128.66666666666666|
+-----+-----+
```

only showing top 10 rows

In [100]: *# Feature 4: number of sessions*

```
user_session_count = df_churn.select("userId", "sessionId").dropDuplicat
es().groupby("userId").count()
user_session_count = user_session_count.withColumnRenamed("count", "sess
ionCount")
```

```
user_session_count.show(5)
```

```
+-----+-----+
|userId|sessionCount|
+-----+-----+
| 100010|           7|
| 200002|           6|
|     125|           1|
|      51|          10|
|     124|          29|
+-----+-----+
```

only showing top 5 rows

```
In [87]: # Feature 5: frequency of use of pages

# get all the type of page
page_list = [(row['page']) for row in df_churn.select("page").dropDuplicates().collect()]

# must remove the column which will cause data leakage
page_list.remove("Cancel")
page_list.remove("Cancellation Confirmation")

# caculate the total page each user view
user_page_view_count = df_churn.groupby("userId").count()
user_page_view_count = user_page_view_count.withColumnRenamed("count", "pageCount")

for page in page_list:
    col_name = "count" + page.replace(" ", "")
    view_count = df_churn.filter(df_churn.page==page).groupby("userId").count()
    view_count = view_count.withColumnRenamed("count", col_name).withColumnRenamed("userId", "userIdTemp")
    user_page_view_count = user_page_view_count.join(view_count, user_page_view_count.userId==view_count.userIdTemp, "left").drop("userIdTemp")
    user_page_view_count = user_page_view_count.sort("userId")
    user_page_view_count = user_page_view_count.fillna(0)

col_list = user_page_view_count.columns
col_list.remove("userId")
col_list.remove("pageCount")
freq_sql = "select userId"
for col in col_list:
    col_name = col.replace("count", "freq")
    sql_str = ", (" + col + "/(pageCount/100)) as " + col_name
    freq_sql = freq_sql + sql_str
freq_sql = freq_sql + " from user_page_view_count"

user_page_view_count.createOrReplaceTempView("user_page_view_count")
col_list = user_page_view_count.columns
col_list.remove("userId")
col_list.remove("pageCount")
freq_sql = "select userId"
for col in col_list:
    col_name = col.replace("count", "freq")
    sql_str = ", (" + col + "/(pageCount/100)) as " + col_name
    freq_sql = freq_sql + sql_str
freq_sql = freq_sql + " from user_page_view_count"

user_page_view_freq = spark.sql(freq_sql)
```

```
In [89]: user_page_view_freq.show(20)
```

```

+-----+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+-----+-----+
|userId| freqSubmitDowngrade|      freqThumbsDown|      freqHome|
freqDowngrade|      freqRollAdvert|      freqLogout|      freqSaveSetting
s|      freqAbout|      freqSettings| freqAddtoPlaylist|      freqA
ddFriend|      freqNextSong|      freqThumbsUp|      freqHelp|
freqUpgrade|      freqError|      freqSubmitUpgrade|
+-----+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+-----+-----+
|      10|      0.0|0.5031446540880503|3.7735849056603774| 0.8
805031446540881|0.12578616352201258|1.3836477987421383|0.12578616352201
258|0.25157232704402516| 0.8805031446540881|1.1320754716981132|1.509433
9622641508| 84.65408805031447| 4.654088050314465|0.12578616352201258|
0.0|      0.0|      0.0|
|      100|0.031113876789047916|0.8400746733042938| 3.266957062850031| 0.9
334163036714375| 0.7778469197261979| 1.088985687616677| 0.1555693839452
396|0.37336652146857496|0.34225264467952704| 1.897946484131923|1.524579
9626633478| 83.44741754822651| 4.604853764779091| 0.5600497822028625|0.
031113876789047916|0.09334163036714374|0.031113876789047916|
|100001|      0.0| 1.06951871657754| 5.88235294117647|
0.0| 7.4866310160427805|3.7433155080213902|      0.0|
0.0| 0.53475935828877|1.6042780748663101| 1.06951871657754| 71.12299
465240642| 4.27807486631016| 0.53475935828877| 1.06951871657754|
0.53475935828877|      0.0|
|100002|      0.0|      0.0|2.7522935779816513| 0.9
174311926605504| 1.3761467889908257|0.4587155963302752|
0.0|      0.0|      0.0| 2.293577981651376|0.458715
5963302752| 89.44954128440367| 2.293577981651376|      0.0|
0.0|      0.0|      0.0|
|100003|      0.0|      0.0| 8.974358974358974|
0.0| 11.538461538461538| 3.846153846153846|      0.0|
0.0|      0.0| 2.564102564102564|      0.0| 65.38461
538461539| 3.846153846153846| 1.282051282051282|      0.0|
0.0|      0.0|
|100004| 0.1606425702811245|0.8835341365461847| 5.301204819277109| 0.8
032128514056225| 6.907630522088354|1.5261044176706828| 0.1606425702811
245|      0.0| 0.8835341365461847|1.8473895582329318|1.526104
4176706828| 75.66265060240964| 2.811244979919679|0.48192771084337355|
0.642570281124498| 0.1606425702811245| 0.24096385542168677|
|100005|      0.0|1.3888888888888888| 6.944444444444444|
0.0| 8.333333333333332|1.3888888888888888|      0.0|
0.0| 0.9259259259259258|1.3888888888888888|1.3888888888888888| 71.29629
629629629|3.2407407407407405| 0.9259259259259258| 1.8518518518518516|
0.0|      0.0|
|100006|      0.0| 4.545454545454546| 4.545454545454546|
0.0| 6.818181818181818| 2.272727272727273|      0.0|
0.0| 2.272727272727273| 2.272727272727273| 9.090909090909092| 59.09090
909090909| 4.545454545454546|      0.0|      0.0|
0.0|      0.0|
|100007|      0.0|1.1538461538461537| 3.846153846153846| 1.
346153846153846| 0.9615384615384615|0.9615384615384615| 0.1923076923076

```

923| 0.0| 0.5769230769230769|1.7307692307692306| 3.26923  
0769230769| 81.34615384615384|3.6538461538461537| 0.5769230769230769|  
0.0| 0.0| 0.0|  
|100008| 0.10638297872340426|0.6382978723404255| 2.553191489361702| 1.0  
638297872340425| 2.127659574468085|0.7446808510638298|  
0.0| 0.3191489361702127| 0.3191489361702127|3.1914893617021276|1.808510  
6382978722| 82.12765957446808|3.9361702127659575| 0.6382978723404255|  
0.425531914893617| 0.0| 0.0|  
|100009| 0.14903129657228018|1.1922503725782414| 3.427719821162444| 0.7  
451564828614009| 6.259314456035767|1.9374068554396424|0.14903129657228  
018|0.14903129657228018| 0.5961251862891207|1.7883755588673622|1.043219  
0760059612| 77.19821162444113| 3.427719821162444| 0.8941877794336811|  
0.5961251862891207| 0.0| 0.14903129657228018|  
|100010| 0.0|1.3123359580052494|2.8871391076115485|  
0.0| 13.648293963254593|1.3123359580052494| 0.0|0.262467  
19160104987| 0.0| 1.837270341207349|1.0498687664041995|  
72.17847769028872|4.4619422572178475| 0.5249343832020997| 0.5249343832  
020997| 0.0| 0.0|  
|100011| 0.0|4.3478260869565215|17.391304347826086|  
0.0| 8.695652173913043| 0.0| 0.0|  
0.0| 4.3478260869565215| 8.695652173913043| 0.0|47.826086  
956521735| 0.0| 0.0| 0.0|  
0.0| 0.0|  
|100012| 0.1666666666666666| 1.5| 4.5| 0.6  
6666666666666666| 6.333333333333333| 1.0|  
0.0|0.1666666666666666| 0.0| 2.0|0.333333  
3333333333| 79.33333333333333| 3.0| 0.3333333333333333|  
0.1666666666666666| 0.0| 0.1666666666666666|  
|100013| 0.0|1.0775862068965518| 3.807471264367816| 0.9  
339080459770115| 2.8017241379310347|0.9339080459770115|0.21551724137931  
036|0.14367816091954022| 0.5028735632183908|2.2270114942528734|2.011494  
2528735633| 81.25|2.8017241379310347| 0.646551724137931|  
0.28735632183908044|0.14367816091954022| 0.07183908045977011|  
|100014| 0.0|0.9677419354838709| 2.258064516129032| 0.9  
677419354838709| 0.6451612903225806|0.9677419354838709|  
0.0| 0.0| 0.3225806451612903| 2.258064516129032|1.935483  
8709677418| 82.90322580645162| 5.483870967741935| 0.6451612903225806|  
0.0| 0.0| 0.0|  
|100015| 0.09523809523809523|0.7619047619047619| 4.476190476190476|0.47  
619047619047616| 6.571428571428571|1.7142857142857142|0.09523809523809  
523|0.38095238095238093| 0.6666666666666666|2.0952380952380953|1.333333  
3333333333| 76.19047619047619|3.3333333333333335| 0.5714285714285714|  
0.7619047619047619|0.09523809523809523| 0.19047619047619047|  
|100016| 0.15673981191222572|0.7836990595611285| 3.29153605015674| 0.9  
404388714733543| 2.5078369905956115|1.4106583072100314|  
0.0|0.15673981191222572|0.15673981191222572|0.9404388714733543|2.037617  
5548589344| 83.07210031347962| 3.918495297805643| 0.4702194357366771|  
0.0|0.15673981191222572| 0.0|  
|100017| 0.0|1.3333333333333333| 4.0|  
0.0| 18.666666666666668| 0.0| 0.0|  
0.0| 0.0|1.3333333333333333| 0.0| 69.33333  
3333333333|2.6666666666666665| 0.0| 0.0|  
0.0| 0.0|  
|100018| 0.15527950310559005|0.6987577639751552|3.3385093167701863| 0.6  
987577639751552| 6.211180124223602|1.3198757763975155| 0.2329192546583  
851|0.07763975155279502| 0.6987577639751552|2.4068322981366457|1.785714  
2857142856| 77.79503105590062| 3.571428571428571| 0.5434782608695652|

```

0.3105590062111801| 0.0| 0.15527950310559005|
+-----+-----+-----+-----+-----+-----+
|-----+-----+-----+-----+-----+-----+
|-----+-----+-----+-----+-----+-----+
|-----+-----+-----+-----+-----+-----+
|-----+-----+-----+-----+-----+-----+
|-----+-----+-----+-----+-----+-----+
only showing top 20 rows

```

In [90]: *# Feature 6: how many singers have the user heard*

```

user_artist_count = df_churn.filter(df_churn.page=="NextSong").select("u
serId", "artist").dropDuplicates().groupBy("userId").count()
user_artist_count = user_artist_count.withColumnRenamed("count", "aritst
Count")

user_artist_count.show(10)

```

```

+-----+-----+
|userId|aritstCount|
+-----+-----+
|100010|      252|
|200002|      339|
|   125|         8|
|    51|     1385|
|   124|     2232|
|     7|      142|
|    15|     1302|
|    54|     1744|
|   155|      643|
|   132|     1299|
+-----+-----+
only showing top 10 rows

```

In [91]: *#. Label: churn*

```
user_churn = df_churn.select("userId", "churn").dropDuplicates()
user_churn = user_churn.select("userId", user_churn.churn.cast("int"))

user_churn.show(10)
```

```
+-----+-----+
|userId|churn|
+-----+-----+
|    19|    0|
|100005|    1|
|200007|    0|
|300007|    1|
|    50|    0|
|200002|    0|
|    30|    0|
|     8|    0|
|100011|    1|
|100012|    1|
+-----+-----+
```

only showing top 10 rows

In [101]: *# Putting all the features into a dataframe*

```
# put all the features dataframe into a list
features_list = []
features_list.append(user_reg_days)
features_list.append(user_session_songs_avg)
features_list.append(user_session_count)
features_list.append(user_gender)
features_list.append(user_page_view_freq)
features_list.append(user_artist_count)
features_list.append(user_churn)
```

In [102]: *# prepare the final dataframe to join all the other features*

```
df_final = df_churn.select("userId").dropDuplicates()

def features_merge(df1, df2):
    """
    This function is used to merge the feature using left join
    input: two data frame to be merged
    output: merged dataframe
    """
    df2 = df2.withColumnRenamed("userId", "userIdTemp")
    df = df1.join(df2, df1.userId == df2.userIdTemp, "left").drop("userIdTemp")
    return df
```



```
In [103]: # use function to merge the features in the list
          for feature in features_list:
              df_final = features_merge(df_final, feature)

          # sort and view the final dataframe
          df_final = df_final.sort("userId")
          df_final.persist()
          df_final.show(5)
```

```

+-----+-----+-----+-----+-----+-----+
|userId|      regDay|  avgSessionSongs|sessionCount|gender| freq
SubmitDowngrade|      freqThumbsDown|      freqHome|      freqDowngrad
e|      freqRollAdvert|      freqLogout|      freqSaveSettings|
freqAbout|      freqSettings| freqAddtoPlaylist|      freqAddFriend|
freqNextSong|      freqThumbsUp|      freqHelp|      freqUpgrade
|      freqError|      freqSubmitUpgrade|aritstCount|churn|
+-----+-----+-----+-----+-----+-----+
|      10| 51.76265046296297|112.16666666666667|      6|      0|
0.0|0.5031446540880503|3.7735849056603774|0.8805031446540881|0.12578616
352201258|1.3836477987421383|0.12578616352201258|0.25157232704402516|
0.8805031446540881|1.1320754716981132|1.5094339622641508|84.65408805031
447|4.654088050314465|0.12578616352201258|      0.0|
0.0|      0.0|      565|      0|
|      100| 64.87377314814815| 78.88235294117646|      35|      0|0.031
113876789047916|0.8400746733042938| 3.266957062850031|0.933416303671437
5| 0.7778469197261979| 1.088985687616677| 0.1555693839452396|0.37336652
146857496|0.34225264467952704| 1.897946484131923|1.5245799626633478|83.
44741754822651|4.604853764779091| 0.5600497822028625|0.0311138767890479
16|0.09334163036714374|0.031113876789047916|      1705|      0|
|100001| 44.80021990740741|      33.25|      4|      1|
0.0| 1.06951871657754| 5.88235294117647|      0.0| 7.4866310
160427805|3.7433155080213902|      0.0|      0.0|
0.53475935828877|1.6042780748663101| 1.06951871657754|71.1229946524064
2| 4.27807486631016| 0.53475935828877| 1.06951871657754| 0.53475
935828877|      0.0|      125|      1|
|100002|160.47207175925925|      48.75|      4|      1|
0.0|      0.0|2.7522935779816513|0.9174311926605504| 1.3761467
889908257|0.4587155963302752|      0.0|      0.0|
0.0| 2.293577981651376|0.4587155963302752|89.44954128440367|2.293577981
651376|      0.0|      0.0|      0.0|
0.0|      184|      0|
|100003|22.748113425925926|      25.5|      2|      1|
0.0|      0.0| 8.974358974358974|      0.0| 11.538461
538461538| 3.846153846153846|      0.0|      0.0|
0.0| 2.564102564102564|      0.0|65.38461538461539|3.846153846
153846| 1.282051282051282|      0.0|      0.0|
0.0|      50|      1|
+-----+-----+-----+-----+-----+-----+
only showing top 5 rows

```

In [ ]:

## Modeling

Split the full dataset into train, test, and validation sets. Test out several of the machine learning methods you learned. Evaluate the accuracy of the various models, tuning parameters as necessary. Determine your winning model based on test accuracy and report results on the validation set. Since the churned users are a fairly small subset, I suggest using F1 score as the metric to optimize.

In [104]: *# let's check how much data do we have.*

```
df_final.groupby("churn").count().show()
```

```
+-----+-----+
| churn | count |
+-----+-----+
|      1 |     52 |
|      0 |    173 |
+-----+-----+
```

In [105]: output = "final\_data.csv"

```
df_final.write.save(output, format="csv", header=True)
```

```
df_final = spark.read.csv(output, header=True)
df_final.persist()
```

Out[105]: DataFrame[userId: string, regDay: string, avgSessionSongs: string, sessionCount: string, gender: string, freqSubmitDowngrade: string, freqThumbsDown: string, freqHome: string, freqDowngrade: string, freqRollAdvert: string, freqLogout: string, freqSaveSettings: string, freqAbout: string, freqSettings: string, freqAddtoPlaylist: string, freqAddFriend: string, freqNextSong: string, freqThumbsUp: string, freqHelp: string, freqUpgrade: string, freqError: string, freqSubmitUpgrade: string, artistCount: string, churn: string]

In [106]: *# Lets convert all the features to numeric.*

```
num_features_list = df_final.columns[1:]
for f in num_features_list:
    f_name = f + "Num"
    df_final = df_final.withColumn(f_name, df_final[f].cast("float"))
    df_final = df_final.drop(f)
```

In [107]: *# Lets train the features into a vector*

```
assembler = VectorAssembler(inputCols=df_final.columns[1:-1], outputCol="NumFeatures")
data = assembler.transform(df_final)

scaler = StandardScaler(inputCol="NumFeatures", outputCol="ScaledNumFeatures", withStd=True)
scalerModel = scaler.fit(data)
data = scalerModel.transform(data)

data = data.select(data.churnNum.alias("label"), data.ScaledNumFeatures.alias("features"))

train, validation = data.randomSplit([0.9, 0.1], seed=42)
train = train.cache()
```

In [108]:

```
lr = LogisticRegression()
paramGrid = ParamGridBuilder() \
    .addGrid(lr.elasticNetParam,[0.0, 0.1, 0.5, 1.0]) \
    .addGrid(lr.regParam,[0.0, 0.05, 0.1]) \
    .build()

crossval = CrossValidator(estimator=lr,
                          estimatorParamMaps=paramGrid,
                          evaluator=MulticlassClassificationEvaluator(),
                          numFolds=3)

cvModel_lr = crossval.fit(train)

cvModel_lr.save('cvModel_lr.model')

cvModel_lr.avgMetrics
```

Out[108]: [0.8005825382361194,  
0.7373409099969138,  
0.7331116591529545,  
0.8005825382361194,  
0.7288125875568267,  
0.7284167395575718,  
0.8005825382361194,  
0.7372280827767593,  
0.6951150985956462,  
0.8005825382361194,  
0.7119161587284445,  
0.6935331299204026]

```
In [109]: dt = DecisionTreeClassifier()
paramGrid = ParamGridBuilder() \
    .addGrid(dt.impurity,['entropy', 'gini']) \
    .addGrid(dt.maxDepth,[2, 3, 4, 5, 6, 7, 8]) \
    .build()

crossval_dt = CrossValidator(estimator=dt,
                             estimatorParamMaps=paramGrid,
                             evaluator=MulticlassClassificationEvaluator(),
                             numFolds=3)

cvModel_dt = crossval_dt.fit(train)

cvModel_dt.save('cvModel_dt.model')

cvModel_dt.avgMetrics
```

```
Out[109]: [0.809391103951177,
0.7501008119218004,
0.7243218701425422,
0.7125347861131004,
0.7199490636115637,
0.7336895542612214,
0.7310787550540219,
0.8245202237499156,
0.7383282194663868,
0.7446536981714925,
0.752821690006628,
0.753965808893903,
0.7417978017336555,
0.7417978017336555]
```

```
In [112]: gbt = GBTClassifier()

paramGrid = ParamGridBuilder() \
    .addGrid(gbt.maxIter,[3, 10, 20]) \
    .addGrid(gbt.maxDepth,[2, 4, 6, 8]) \
    .build()

crossval_gbt = CrossValidator(estimator=gbt,
                              estimatorParamMaps=paramGrid,
                              evaluator=MulticlassClassificationEvaluator(),
                              numFolds=3)

cvModel_gbt = crossval_gbt.fit(train)

cvModel_gbt.save('cvModel_gbt.model')

cvModel_gbt.avgMetrics
```

```
Out[112]: [0.823855068054212,
0.7378619062152055,
0.753965808893903,
0.7417978017336555,
0.7948841221807386,
0.7418057214032604,
0.753965808893903,
0.7417978017336555,
0.7855099580040472,
0.7338696709530763,
0.753965808893903,
0.7417978017336555]
```

```
In [113]: stratified_train = train.sampleBy('label', fractions={0: 99/349, 1: 1.0
}).cache()
stratified_train.groupby("label").count().show()
```

```
+-----+-----+
|label|count|
+-----+-----+
|  1.0|   44|
|  0.0|   40|
+-----+-----+
```

```
In [114]: lrs = LogisticRegression()
paramGrid = ParamGridBuilder() \
    .addGrid(lrs.elasticNetParam,[0.0, 0.1, 0.5, 1.0]) \
    .addGrid(lrs.regParam,[0.0, 0.05, 0.1]) \
    .build()

crossval_lrs = CrossValidator(estimator=lrs,
                              estimatorParamMaps=paramGrid,
                              evaluator=MulticlassClassificationEvaluator(),
                              numFolds=3)
cvModel_lrs = crossval_lrs.fit(stratified_train)
cvModel_lrs.avgMetrics
```

```
Out[114]: [0.584848484848485,
0.5685105018438352,
0.5588924777779267,
0.584848484848485,
0.5362999712979521,
0.5616072777198097,
0.584848484848485,
0.5717316878442198,
0.5106045607120876,
0.584848484848485,
0.5823046066275844,
0.5174310801761782]
```

```
In [115]: cvModel_lrs.save('cvModel_lrs.model')
```

```
In [116]: dts = DecisionTreeClassifier()
paramGrid = ParamGridBuilder() \
    .addGrid(dts.impurity,['entropy', 'gini']) \
    .addGrid(dts.maxDepth,[2, 3, 4, 5, 6, 7, 8]) \
    .build()
crossval_dts = CrossValidator(estimator=dts,
                              estimatorParamMaps=paramGrid,
                              evaluator=MulticlassClassificationEvaluator(),
                              numFolds=3)
cvModel_dts = crossval_dts.fit(stratified_train)
cvModel_dts.avgMetrics
```

```
Out[116]: [0.5995466555416493,
0.6720397283443197,
0.5931161073096557,
0.5588591026091027,
0.6150359347281962,
0.6061778499278501,
0.6061778499278501,
0.5835235065245915,
0.6264957264957265,
0.5459702496983199,
0.5365808014422584,
0.5365808014422584,
0.5365808014422584,
0.5365808014422584]
```

```
In [117]: cvModel_dts.save('cvModel_dts.model')
```

```
In [118]: gbts = GBTCClassifier()
paramGrid = ParamGridBuilder() \
    .addGrid(gbts.maxIter,[3, 10, 20]) \
    .addGrid(gbts.maxDepth,[2, 4, 6, 8]) \
    .build()
crossval_gbts = CrossValidator(estimator=gbts,
                               estimatorParamMaps=paramGrid,
                               evaluator=MulticlassClassificationEvaluator(),
                               numFolds=3)
cvModel_gbts = crossval_gbts.fit(stratified_train)
cvModel_gbts.avgMetrics
```

```
Out[118]: [0.6027873524138243,
0.5499410164960218,
0.5365808014422584,
0.5365808014422584,
0.6568068172187166,
0.5945591087526572,
0.5365808014422584,
0.5365808014422584,
0.66054381348499,
0.6262132928799596,
0.5365808014422584,
0.5365808014422584]
```

```
In [119]: cvModel_gbts.save('cvModel_gbts.model')
```

```
In [120]: # Use the validate data to evaluate the best model through F1 score.
```

```
results = cvModel_lr.transform(validation)

tp = results.filter("label = 1 and prediction = 1").count()
fp = results.filter("label = 0 and prediction = 1").count()
fn = results.filter("label = 1 and prediction = 0").count()
precision = tp / (tp + fp)
recall = tp / (tp + fn)
f1 = 2*precision*recall / (precision+recall)
print(precision)
print(recall)
print(f1)
```

```
1.0
0.75
0.8571428571428571
```

```
In [121]: results = cvModel_lrs.transform(validation)
```



```
In [122]: tp = results.filter("label = 1 and prediction = 1").count()
fp = results.filter("label = 0 and prediction = 1").count()
fn = results.filter("label = 1 and prediction = 0").count()
precision = tp / (tp + fp)
recall = tp / (tp + fn)
f1 = 2*precision*recall / (precision+recall)
print(precision)
print(recall)
print(f1)
```

```
0.6363636363636364
0.875
0.7368421052631579
```

```
In [123]: cvModel_lrs.bestModel.coefficients
```

```
Out[123]: DenseVector([4.8203, 11.436, -16.637, 5.6112, -13.3528, -107.9358, -27
0.1181, -87.8711, -572.666, -119.6136, -44.9646, -27.1968, -233.8042, -
172.517, -235.4224, -953.0062, -285.0545, -59.4815, -55.8425, -30.8607,
-27.3215, 21.452])
```

```
In [124]: df_final.columns
```

```
Out[124]: ['userId',
'regDayNum',
'avgSessionSongsNum',
'sessionCountNum',
'genderNum',
'freqSubmitDowngradeNum',
'freqThumbsDownNum',
'freqHomeNum',
'freqDowngradeNum',
'freqRollAdvertNum',
'freqLogoutNum',
'freqSaveSettingsNum',
'freqAboutNum',
'freqSettingsNum',
'freqAddtoPlaylistNum',
'freqAddFriendNum',
'freqNextSongNum',
'freqThumbsUpNum',
'freqHelpNum',
'freqUpgradeNum',
'freqErrorNum',
'freqSubmitUpgradeNum',
'aritstCountNum',
'churnNum']
```

```
In [ ]:
```

```
In [ ]:
```

## Final Steps

Clean up your code, adding comments and renaming variables to make the code easier to read and maintain. Refer to the Spark Project Overview page and Data Scientist Capstone Project Rubric to make sure you are including all components of the capstone project and meet all expectations. Remember, this includes thorough documentation in a README file in a Github repository, as well as a web app or blog post.

In [ ]: